Financial Market Data Pipeline

Data Preparation Phase – Project Update

Establishing clean, validated, and analytics-ready datasets for corporate stakeholders.
Prepared for executive review and data strategy alignment.

Prepared by: Nandani Raj

MS Business Analytics – University of New Haven

March 10, 2026

Project Objective

Standardize Multi-Source Data

Build a scalable financial data pipeline to integrate and standardize multi-source market data.

Execute ETL Processes

Transform raw data into structured, analytics-ready formats via ETL **(Extract, Transform, Load – moving and preparing data for analysis)**.

Power Financial Analytics

Enable dynamic dashboards, robust investment insights, and repeatable financial performance tracking.

Ensure Traceability & Governance

Maintain strict traceability, reproducibility, and structural governance across all data layers.

Drive Business Value

Deliver a highly reliable data foundation specifically designed for executive decision-support analytics.

Business Problem We Solve

Inconsistent Multi-Source Data

Financial data arrives from various external sources with completely different schemas and data structures.

Complex & Incompatible Formats

Raw datasets often mix metadata with prices and use wide formats, creating severe structural barriers to dataset integration.

Inefficient Resource Allocation

Data analysts and scientists spend excessive, high-value time manually cleaning data instead of generating strategic insights.

Automated Pipeline Standardization

Our solution introduces an automated architecture that enforces structured schemas and strictly unifies disjointed sources.

High-Impact Business Outcomes

This foundation directly enables significantly faster time-to-insight and trustworthy comparative analytics.

Data Sources Overview

1. Stock Prices Dataset

Historical daily trading data for five major technology companies (AAPL, MSFT, AMZN, NVDA, TSLA). Supports stock-level trend and performance analysis.

Coverage: 12,581 records

2. S&P 500 Market Index (^GSPC)

Daily performance representing the overall U.S. stock market.

Coverage: 2,517 records

Serves as a **Market Benchmark: A reference standard to assess relative stock performance.*

3. S&P 500 Company Metadata

Company-level contextual information. Enriches financial analysis by enabling sector and industry classifications.

Coverage: 503 company records

Data Structure & Variables

Time-Series Data

Open | High | Low | Close | Volume

Date | **Ticker**

Business Purpose: The OHLCV structure is the industry standard format used for financial market analysis and technical charting.

Company Metadata

Ticker | Name | Sector | Industry

HQ Location | Date Added | CIK / Founded

Business Purpose: Provides crucial dimensions for aggregating data, enabling cross-industry and geographic performance comparisons.

Integration Keys: Ticker Symbol & Date

Data Extraction Process

Automated retrieval of financial market data with an emphasis on preserving source integrity. Our pipeline establishes a structured ingestion flow before any data cleaning occurs.

1. API Extraction

Python scripts using the **yFinance API** automatically pull historical stock prices and market index data with reproducible parameters.

2. Raw Data Layer

Source data is loaded into a dedicated raw data landing zone.

(Stored exactly in its original format)

3. Immutable Storage

No modifications or cleaning occur at this stage. Versioning and extraction timestamps are rigorously preserved for safety.

4. Traceability

Ensures full auditability, provides the capability to re-run pipelines without API limits, and guarantees consistent cleaning inputs.

Architecture Benefit: Maintaining a raw data layer ensures the pipeline is traceable, auditable, and easily re-runnable.

Data Cleaning – Stock Prices Dataset

Original Dataset Assessment

Evaluated initial shape of 12,581 rows × 27 columns, noting multiple tickers embedded as wide column groups.

Metadata Row Removal

Removed non-trading header metadata rows to ensure the dataset strictly contains valid trading observations.

Data Type Standardization

Converted Date fields to strictly datetime format and financial metrics (OHLCV) to numeric data types.

Wide-to-Long Transformation (Melt)

Restructured data into a long format (one standardized row per Ticker-Date-Variable) to ensure analytical compatibility.

Business Implications

Scalability & Integration Ready

The long format structure is a foundational requirement for relational databases and analytical workflows.

By eliminating wide, unscalable columns, the dataset is now instantly join-ready with other market feeds, unlocking reliable cross-sectional and time-series analytics.

Data Cleaning – S&P 500 Market Index

Initial Dataset Assessment

Evaluated daily OHLCV data for the index with an initial structural shape of 2,517 rows × 6 columns.

Metadata Row Removal

Identified and removed the non-trading header row (containing the ticker ^GSPC) to isolate valid market records.

Format & Type Standardization

Standardized column names, enforced strict numeric types for financial variables, and converted Date to datetime.

Financial Logic Validation

Validated underlying numerical relationships, verifying logical consistency across records (e.g., Low ≤ Open/Close ≤ High).

Market Benchmark

Why Index Data Matters

A market benchmark is a standard reference index (like the S&P 500) used to represent the overall performance of the broader market.

By securing a clean index dataset, analysts can definitively answer whether individual stock movements are driven by company-specific events or wider macroeconomic trends.

Data Cleaning – Company Metadata

Dataset Assessment

Evaluated 503 S&P 500 company records, confirming completeness across 8 distinct corporate metadata fields.

Standardization & Validation

Standardized column names and rigorously validated ticker symbols to ensure accurate joining with price datasets.

Missing Value Analysis

Inspected dataset structure and confirmed **zero missing values or duplicate records** across all entries.

Key Attributes Verified

Retained critical contextual data including **Sector, Industry, Headquarters, CIK, Date added, and Founded year**.

Business Implications

Sector & Industry Analytics

The metadata dataset is structurally complete and fully ready to enrich raw stock prices with business context.

By merging this metadata, the pipeline enables **sector-level performance comparisons**, industry concentration analysis, and geographic distribution insights.

Data Quality Validation

Data Validation is the systematic process of ensuring data accuracy and structural integrity. Implementing strict validation checks guarantees the final datasets provide a reliable, error-free foundation for downstream financial analytics and dashboarding.

Column Structure Verification	Ensures schema consistency and expected column presence across all datasets, enabling seamless merging without misalignment errors.
Data Type Enforcement	Validates that chronological fields are strict datetime objects and financial metrics (prices, volumes) are precise numeric formats.
Missing Value Detection	Identifies gaps or null records in critical fields to prevent skewed rolling averages or inadvertently dropped observations during integration.
Duplicate Record Detection	Removes redundant daily observations to ensure total trading volumes, volatility metrics, and index performances are not artificially inflated.
Financial Logic Validation	Confirms strict business rules are inherently met (e.g., Daily High ≥ Daily Low, trading volumes ≥ 0) to maintain absolute analytical integrity.

Current Project Status

Data Preparation Milestones

1. Data Extraction Completed

Historical data successfully retrieved for all core stock tickers and the S&P 500 market benchmark.

2. Dataset Cleaning Finalized

Transformations successfully applied across all sources:

- Stock Prices
- S&P 500 Index
- Company Metadata

3. Data Quality Validation Passed

All schema, data type, logic, and missing-value checks passed. Structural integrity is confirmed.

Phase 1: 100% Complete

All datasets are structured, standardized, and analytics-ready.

Artifacts Produced

- **Raw Data Layer**
  Immutable storage of source data.
- **Cleaned Tables**
  Standardized, long-format datasets ready for DB.
- **Validation Logs**
  Documented quality assurance records.

Governance & Controls

Reproducible Scripts

All extraction and cleaning logic is maintained in version-controlled Python scripts, ensuring the pipeline can be re-run automatically.

Documented Assumptions

Business logic, data types, and gap-handling decisions are fully documented.

Next Phase – Integration & Features

Building a unified financial analytics dataset through strategic dataset merging and analytical feature engineering.

1. Context Integration

Join the Stock Prices dataset with Company Metadata using the **Ticker symbol** to enrich records with Sector and Industry classifications.

2. Market Alignment

Align stock performance dates seamlessly with the **S&P 500 index** to baseline performance and compute relative market metrics.

3. Feature Engineering

Calculate key indicators including **daily returns, volatility, rolling averages**, beta vs index, and sector-level aggregates.

4. Target Data Model

Produce a unified, **analytics-friendly schema** optimized for dashboard ingestion, rapid queries, and downstream ML modeling.

Expected Business Impact

This integrated structure ensures stakeholders can generate deeper insights via reliable **cross-sectional** (e.g., sector comparisons) and **time-series** (e.g., historical trends) views, powering robust decision-support analytics.

Business Value Enabled

The integrated, analytics-ready dataset empowers multiple downstream business capabilities, transforming raw feeds into actionable insights.

Investment Analysis

Identify outperforming stocks by comparing individual performance against the S&P 500 market benchmark and sector peers.

Market Insights

Uncover macro trends, analyze volatility patterns, and leverage time-series data to detect regime signals.

Sector Comparisons

Generate sector and industry heatmaps to analyze geographic distribution and contribution to market performance.

BI Enablement

Provide reliable, structured data feeds to power executive financial dashboards and automated reporting pipelines.

Decision Support

Accelerate time-to-insight with consistent analytical metrics, full auditability, and trustworthy foundational data.

Strategic Impact: The pipeline ensures analysts spend less time cleaning data and more time generating high-value investment research.